

System and Device Programming

Lab 04 Exercises – barriers, memory map and threads

Exercise 1

Write a C program using Pthreads to sort the content of a **binary** file including a sequence of random integer numbers, passed as an argument of the command line.

Map the file as a vector in memory.

Implement a threaded quicksort program where the recursive calls to **quicksort** are replaced by threads activations, i.e. sorting is done, in parallel, in different regions of the file.

If the difference between the **right** and **left** indexes is less than a value **size**, given as an argument of the command line, sorting is performed by the standard **quicksort** algorithm.

This is a sequential recursive implementation of the quicksort algorithm.

```
void quicksort (int v[], int left, int right) {
    int i, j, x, tmp;
    if (left >= right)    return;
    x = v[left];
    i = left - 1;
    j = right + 1;
    while (i < j) {
        while (v[--j] > x);
        while (v[++i] < x);
        if (i < j)
            swap (i,j);
    }
    quicksort (v, left, j);
    quicksort (v, j + 1, right);
}

void swap(int i, int j){
    int tmp;

    tmp = v[i];
    v[i] = v[j];
    v[j] = tmp;
}
```

Exercise 2

Implement a sequential program in C that takes a single argument **k** from the command line. The program creates two vectors (**v1** and **v2**) of dimension **k**, and a matrix (**mat**) of dimension **kxk**, which are filled with random numbers in the range **[-0.5 0.5]**, then it performs the product

v1^T * mat * v2, and print the result. This is an example for **k=5**:

```
v1T = [-0.0613   -0.1184   0.2655   0.2952  -0.3131]
mat=[  -0.3424   -0.3581   0.1557   0.2577   0.2060
      0.4706   -0.0782   -0.4643   0.2431   -0.4682
      0.4572   0.4157   0.3491  -0.1078   -0.2231
     -0.0146   0.2922   0.4340   0.1555   -0.4538
      0.3003   0.4595   0.1787  -0.3288   -0.4029]
v2T = [-0.3235   0.1948  -0.1829   0.4502  -0.4656]
```

Result: 0.0194

Perform the product operation in two steps: $\mathbf{v} = \mathbf{mat} * \mathbf{v2}$, which produces a new vector \mathbf{v} , and $\mathbf{result} = \mathbf{v1}^T * \mathbf{v}$

Then, write a concurrent program using threads that performs the same task. The main thread creates the vectors, the matrix, and k threads. Then, it waits the termination of the other threads.

Each thread i performs the product of the i -th row vector of \mathbf{mat} and $\mathbf{v2}$, which produces the i -th element of vector \mathbf{v} .

One of the created threads, the **last** one terminating its product operation, performs the final operation $\mathbf{result} = \mathbf{v1}^T * \mathbf{v}$, and prints the result.